

▼ Train a simple convnet on the Fashion MNIST dataset

```
#pip install -U tensorflow --quiet
```

In this, we will see how to deal with image data and train a convnet for image classification task.

▼ Load the `fashion_mnist` dataset

** Use keras.datasets to load the dataset **


```
#import tensorflow as tf
import keras

from keras.datasets import fashion_mnist
(x_train, y_train), (x_test, y_test) = fashion_mnist.load_data()
```

▼ Find no.of samples are there in training and test datasets

```
import matplotlib.pyplot as plt
import numpy as np

img_num = np.random.randint(0, x_test.shape[0]) #Get a random integer between 0 and number of
plt.imshow(x_test[img_num],cmap='gray') #Show the image from test dataset
plt.suptitle('Number: ' + str(y_test[img_num]))
plt.show()
```



```
from keras.utils import np_utils
```

```
print('--- THE DATA ---')
print('x_train shape:', x_train.shape)
print(x_train.shape[0], 'train samples')
print(x_test.shape[0], 'test samples')
```

```
↳ --- THE DATA ---
x_train shape: (60000, 28, 28)
60000 train samples
10000 test samples
```

▼ Find dimensions of an image in the dataset

```
x_train.shape
```

```
↳ (60000, 28, 28)
```

```
x_test.shape
```

```
↳ (10000, 28, 28)
```

▼ Convert train and test labels to one hot vectors

```
** check keras.utils.to_categorical() **
```

```
y_train = np_utils.to_categorical(y_train, 10)
y_test = np_utils.to_categorical(y_test, 10)
```

Normalize both the train and test image data from 0-255 to 0-1

▼ Reshape the data from 28x28 to 28x28x1 to match input dimensions in Conv2D layer

```
#Both steps done at one step
x_train = x_train.reshape(x_train.shape[0], 28, 28, 1).astype('float32')
x_test = x_test.reshape(x_test.shape[0], 28, 28, 1).astype('float32')
x_train /= 255
x_test /= 255
```

```
x_train.shape
```

```
↳ (60000, 28, 28, 1)
```

▼ Import the necessary layers from keras to build the model

```
from keras.layers import Dense, Activation, Dropout, Flatten, Reshape
from keras.layers import Convolution2D, MaxPooling2D
from keras.models import Sequential

Sequential()
```

↳ <keras.engine.sequential.Sequential at 0x7fed25555ef0>

▼ Build a model

** with 2 Conv layers having 32 3x3 filters in both convolutions with relu activations and flatt
fully connected layers (or Dense Layers) having 128 and 10 neurons with relu and softmax activatio
categorical_crossentropy loss with adam optimizer train the model with early stopping patience=5

```
TRAIN = False
BATCH_SIZE = 32
EPOCHS = 10

# Define model
model = Sequential()

# 1st Conv Layer
model.add(Convolution2D(32, 3, 3, input_shape=(28, 28, 1)))
model.add(Activation('relu'))

# 2nd Conv Layer
model.add(Convolution2D(32, 3, 3))
model.add(Activation('relu'))

# Fully Connected Layer
model.add(Flatten())
model.add(Dense(128))
model.add(Activation('relu'))

# Prediction Layer
model.add(Dense(10))
model.add(Activation('softmax'))

# Loss and Optimizer
model.compile(loss='categorical_crossentropy', optimizer='adam', metrics=['accuracy'])

# Store Training Results
early_stopping = keras.callbacks.EarlyStopping(monitor='val_acc', patience=10, verbose=1,
callback_list = [early_stopping]

# Train the model
model.fit(x_train, y_train, batch_size=BATCH_SIZE, nb_epoch=EPOCHS,
           validation_data=(x_test, y_test), callbacks=callback_list)
```

```

↳ /usr/local/lib/python3.6/dist-packages/ipykernel_launcher.py:4: UserWarning: Update your
    after removing the cwd from sys.path.
/usr/local/lib/python3.6/dist-packages/ipykernel_launcher.py:8: UserWarning: Update your

/usr/local/lib/python3.6/dist-packages/ipykernel_launcher.py:29: UserWarning: The `nb_ep
WARNING:tensorflow:From /usr/local/lib/python3.6/dist-packages/tensorflow_core/python/op
Instructions for updating:
Use tf.where in 2.0, which has the same broadcast rule as np.where
WARNING:tensorflow:From /usr/local/lib/python3.6/dist-packages/keras/backend/tensorflow_
WARNING:tensorflow:From /usr/local/lib/python3.6/dist-packages/keras/backend/tensorflow_
WARNING:tensorflow:From /usr/local/lib/python3.6/dist-packages/keras/backend/tensorflow_
WARNING:tensorflow:From /usr/local/lib/python3.6/dist-packages/keras/backend/tensorflow_
Train on 60000 samples, validate on 10000 samples
Epoch 1/10
WARNING:tensorflow:From /usr/local/lib/python3.6/dist-packages/keras/backend/tensorflow_
WARNING:tensorflow:From /usr/local/lib/python3.6/dist-packages/keras/backend/tensorflow_
WARNING:tensorflow:From /usr/local/lib/python3.6/dist-packages/keras/backend/tensorflow_
WARNING:tensorflow:From /usr/local/lib/python3.6/dist-packages/keras/backend/tensorflow_
WARNING:tensorflow:From /usr/local/lib/python3.6/dist-packages/keras/backend/tensorflow_
60000/60000 [=====] - 28s 470us/step - loss: 0.3693 - acc: 0.86
Epoch 2/10
60000/60000 [=====] - 20s 334us/step - loss: 0.2313 - acc: 0.91
Epoch 3/10
60000/60000 [=====] - 20s 333us/step - loss: 0.1670 - acc: 0.93
Epoch 4/10
60000/60000 [=====] - 19s 325us/step - loss: 0.1161 - acc: 0.95
Epoch 5/10
60000/60000 [=====] - 19s 322us/step - loss: 0.0778 - acc: 0.97
Epoch 6/10
60000/60000 [=====] - 19s 321us/step - loss: 0.0505 - acc: 0.98
Epoch 7/10
60000/60000 [=====] - 19s 320us/step - loss: 0.0381 - acc: 0.98
Epoch 8/10
60000/60000 [=====] - 19s 322us/step - loss: 0.0256 - acc: 0.99
Epoch 9/10
60000/60000 [=====] - 19s 321us/step - loss: 0.0256 - acc: 0.99
Epoch 10/10
60000/60000 [=====] - 19s 320us/step - loss: 0.0182 - acc: 0.99
<keras.callbacks.History at 0x7fed2540fd0>

```

Now, to the above model add max pooling layer of filter size 2x2 and dropout conv layers and run the model

```
loss_and_metrics = model.evaluate(x_test, y_test)
print(loss_and_metrics)
```

```
↳ 10000/10000 [=====] - 1s 91us/step  
[0.49650015604533254, 0.911]
```

BATCH_SIZE = 32

EPOCHS = 10

```
# Define Model
model2 = Sequential()

# 1st Conv Layer
model2.add(Convolution2D(32, 3, 3, input_shape=(28, 28, 1)))
model2.add(Activation('relu'))

# 2nd Conv Layer
model2.add(Convolution2D(32, 3, 3))
model2.add(Activation('relu'))

# Max Pooling
model2.add(MaxPooling2D(pool_size=(2,2)))

# Dropout
model2.add(Dropout(0.25))

# Fully Connected Layer
model2.add(Flatten())
model2.add(Dense(128))
model2.add(Activation('relu'))

# More Dropout
model2.add(Dropout(0.5))

# Prediction Layer
model2.add(Dense(10))
model2.add(Activation('softmax'))

# Loss and Optimizer
model2.compile(loss='categorical_crossentropy', optimizer='adam', metrics=['accuracy'])

# Store Training Results
early_stopping = keras.callbacks.EarlyStopping(monitor='val_acc', patience=7, verbose=1,
callback_list = [early_stopping]

# Train the model
model2.fit(x_train, y_train, batch_size=BATCH_SIZE, nb_epoch=EPOCHS,
validation_data=(x_test, y_test), callbacks=callback_list)
```

```
↳
```

```

WARNING:tensorflow:From /usr/local/lib/python3.6/dist-packages/keras/backend/tensorflow_
WARNING:tensorflow:From /usr/local/lib/python3.6/dist-packages/keras/backend/tensorflow_
WARNING:tensorflow:From /usr/local/lib/python3.6/dist-packages/keras/backend/tensorflow_
Instructions for updating:
Please use `rate` instead of `keep_prob`. Rate should be set to `rate = 1 - keep_prob`.
/usr/local/lib/python3.6/dist-packages/ipykernel_launcher.py:4: UserWarning: Update your
    after removing the cwd from sys.path.
/usr/local/lib/python3.6/dist-packages/ipykernel_launcher.py:8: UserWarning: Update your

/usr/local/lib/python3.6/dist-packages/ipykernel_launcher.py:38: UserWarning: The `nb_ep
Train on 60000 samples, validate on 10000 samples
Epoch 1/10
60000/60000 [=====] - 19s 317us/step - loss: 0.5117 - acc: 0.81
Epoch 2/10
60000/60000 [=====] - 18s 308us/step - loss: 0.3493 - acc: 0.87
Epoch 3/10
60000/60000 [=====] - 18s 308us/step - loss: 0.3037 - acc: 0.89
Epoch 4/10
60000/60000 [=====] - 18s 307us/step - loss: 0.2736 - acc: 0.89
Epoch 5/10
60000/60000 [=====] - 18s 308us/step - loss: 0.2513 - acc: 0.90
Epoch 6/10
60000/60000 [=====] - 18s 307us/step - loss: 0.2334 - acc: 0.91
Epoch 7/10
60000/60000 [=====] - 18s 305us/step - loss: 0.2221 - acc: 0.91
Epoch 8/10
60000/60000 [=====] - 18s 306us/step - loss: 0.2069 - acc: 0.92
Epoch 9/10
60000/60000 [=====] - 18s 306us/step - loss: 0.1999 - acc: 0.92
Epoch 10/10
60000/60000 [=====] - 18s 306us/step - loss: 0.1870 - acc: 0.92
<keras.callbacks.History at 0x7fed20334f98>

```

Now, to the above model, lets add Data Augmentation

- ▼ Import the ImageDataGenerator from keras and fit the training images

```

from keras.preprocessing.image import ImageDataGenerator

# This will do preprocessing and realtime data augmentation:
datagen = ImageDataGenerator(
    featurewise_center=False, # set input mean to 0 over the dataset
    samplewise_center=False, # set each sample mean to 0
    featurewise_std_normalization=False, # divide inputs by std of the dataset
    samplewise_std_normalization=False, # divide each input by its std
    zca_whitening=False, # apply ZCA whitening
    rotation_range=50, # randomly rotate images in the range (degrees, 0 to 180)
    width_shift_range=0.01, # randomly shift images horizontally (fraction of total width)
    https://colab.research.google.com/drive/10QJ6w_cCPwiAc-5VTcovNvRvc-N92qTO#scrollTo=mXug4z234mwQ&printMode=true

```

```
height_shift_range=0.01, # randomly shift images vertically (fraction of total height)
horizontal_flip=False, # randomly flip images
vertical_flip=False) # randomly flip images
```

```
# Prepare the generator
datagen.fit(x_train)
```

```
# Define Model
```

```
model3 = Sequential()
```

```
# 1st Conv Layer
```

```
model3.add(Convolution2D(32, 3, 3, input_shape=(28, 28, 1)))
```

```
model3.add(Activation('relu'))
```

```
# 2nd Conv Layer
```

```
model3.add(Convolution2D(32, 3, 3))
```

```
model3.add(Activation('relu'))
```

```
# Max Pooling
```

```
model3.add(MaxPooling2D(pool_size=(2,2)))
```

```
# Dropout
```

```
model3.add(Dropout(0.15))
```

```
# Fully Connected Layer
```

```
model3.add(Flatten())
```

```
model3.add(Dense(128))
```

```
model3.add(Activation('relu'))
```

```
# More Dropout
```

```
model3.add(Dropout(0.2))
```

```
# Prediction Layer
```

```
model3.add(Dense(10))
```

```
model3.add(Activation('softmax'))
```

```
# Loss and Optimizer
```

```
model3.compile(loss='categorical_crossentropy', optimizer='adam', metrics=['accuracy'])
```

```
# Store Training Results
```

```
early_stopping = keras.callbacks.EarlyStopping(monitor='val_acc', patience=7, verbose=1,
callback_list = [early_stopping]
```

↳ /usr/local/lib/python3.6/dist-packages/ipykernel_launcher.py:4: UserWarning: Update your
after removing the cwd from sys.path.

/usr/local/lib/python3.6/dist-packages/ipykernel_launcher.py:8: UserWarning: Update your

▼ Showing 5 versions of the first image in training dataset using image datagenerator.flow()

```
from matplotlib import pyplot as plt
gen = datagen.flow(x_train[0:1], batch_size=1)
for i in range(1, 6):
    plt.subplot(1,5,i)
    plt.axis("off")
    plt.imshow(gen.next().squeeze(), cmap='gray')
    plt.plot()
plt.show()
```



▼ Run the above model using fit_generator()

```
model3.fit_generator(datagen.flow(x_train, y_train,batch_size=32),
                     samples_per_epoch=x_train.shape[0],
                     nb_epoch=10,
                     validation_data=(x_test, y_test), callbacks=callback_list)
```

```
↳ /usr/local/lib/python3.6/dist-packages/ipykernel_launcher.py:4: UserWarning: The semantic of os.getcwd() has changed after removing the cwd from sys.path.
/usr/local/lib/python3.6/dist-packages/ipykernel_launcher.py:4: UserWarning: Update your __file__ attribute after removing the cwd from sys.path.
Epoch 1/10
1875/1875 [=====] - 34s 18ms/step - loss: 0.6573 - acc: 0.7595
Epoch 2/10
1875/1875 [=====] - 33s 18ms/step - loss: 0.4641 - acc: 0.8311
Epoch 3/10
1875/1875 [=====] - 33s 18ms/step - loss: 0.4145 - acc: 0.8481
Epoch 4/10
1875/1875 [=====] - 33s 18ms/step - loss: 0.3827 - acc: 0.8579
Epoch 5/10
1875/1875 [=====] - 33s 18ms/step - loss: 0.3625 - acc: 0.8662
Epoch 6/10
1875/1875 [=====] - 33s 18ms/step - loss: 0.3415 - acc: 0.8746
Epoch 7/10
1875/1875 [=====] - 33s 18ms/step - loss: 0.3313 - acc: 0.8781
Epoch 8/10
1875/1875 [=====] - 33s 18ms/step - loss: 0.3227 - acc: 0.8819
Epoch 9/10
1875/1875 [=====] - 33s 18ms/step - loss: 0.3132 - acc: 0.8841
Epoch 10/10
1875/1875 [=====] - 33s 18ms/step - loss: 0.3055 - acc: 0.8879
<keras.callbacks.History at 0x7fed253c42b0>
```

▼ Report the final train and validation accuracy

```
loss_and_metrics = model3.evaluate(x_train, y_train)
print(loss_and_metrics)

⇒ 60000/60000 [=====] - 5s 87us/step
[0.24732203457752863, 0.9106666666666666]
```

▼ DATA AUGMENTATION ON CIFAR10 DATASET

One of the best ways to improve the performance of a Deep Learning model is to add more data to the instances from the wild that are representative of the distinction task, we want to develop a set of methods to do this. There are many ways to augment existing datasets and produce more robust models. In the image recognition task, we can use data augmentation to fully utilize the full power of the convolutional neural network, which is able to capture translational invariance. This task is challenging because image recognition is such a difficult task in the first place. You want the dataset to be representative of different lighting conditions, different object orientations, different lightings, and miscellaneous distortions that are of interest to the vision task.

▼ Import necessary libraries for data augmentation

```
from keras.preprocessing.image import ImageDataGenerator
```

▼ Load CIFAR10 dataset

```
from keras.datasets import cifar10

(x_train_n, y_train_n), (x_test_n, y_test_n) = cifar10.load_data()

⇒ Downloading data from https://www.cs.toronto.edu/~kriz/cifar-10-python.tar.gz
170500096/170498071 [=====] - 11s 0us/step

y_train_n = np_utils.to_categorical(y_train_n, 10)
y_test_n = np_utils.to_categorical(y_test_n, 10)

print('--- THE DATA ---')
print('x_train shape:', x_train_n.shape)
print(x_train_n.shape[0], 'train samples')
print(x_test_n.shape[0], 'test samples')

⇒ --- THE DATA ---
x_train shape: (50000, 32, 32, 3)
50000 train samples
10000 test samples
```

>Create a data_gen funtion to generator with image rotation,shifting image horizontally random flip horizontally.

```
datagen_n = ImageDataGenerator(
    featurewise_center=False, # set input mean to 0 over the dataset
    samplewise_center=False, # set each sample mean to 0
    featurewise_std_normalization=False, # divide inputs by std of the dataset
    samplewise_std_normalization=False, # divide each input by its std
    zca_whitening=False, # apply ZCA whitening
    rotation_range=60, # randomly rotate images in the range (degrees, 0 to 180)
    width_shift_range=0.01, # randomly shift images horizontally (fraction of total width)
    height_shift_range=0.01, # randomly shift images vertically (fraction of total height)
    horizontal_flip=True, # randomly flip images
    vertical_flip=True) # randomly flip images
```

```
x_train_n = x_train_n.astype('float32')
x_test_n = x_test_n.astype('float32')
```

```
x_train_n /= 255
x_test_n /= 255
```

▼ Prepare/fit the generator.

```
# Prepare the generator
datagen_n.fit(x_train_n)
```

▼ Generate 5 images for 1 of the image of CIFAR10 train dataset.

```
gen = datagen.flow(x_train_n[:1], batch_size=1)
for i in range(1, 6):
    plt.subplot(1,5,i)
    plt.axis("off")
    plt.imshow(gen.next().squeeze())
    plt.plot()
```



